

Building a E-commerce Data Warehouse Using Hive

By Keyur Rajput

Step 1: Setup

- 1.1. Log in to your Cloudera environment.
- 1.2. Start the Hadoop cluster and Hive service:

```
sudo service hadoop-hdfs-datanode start
sudo service hadoop-hdfs-namenode start
sudo service hadoop-yarn-resourcemanager start
sudo service hadoop-yarn-nodemanager start
```

Step 2: Data Ingestion

- 2.1. Begin by downloading a sample e-commerce dataset from Kaggle. You can find a suitable dataset at - <https://www.kaggle.com/datasets/carrie1/ecommerce-data>
- 2.2. Upload the dataset to HDFS using the following commands:

```
hadoop fs -mkdir /user/hive/warehouse/ecommerce
hadoop fs -put /home/cloudera/Desktop/ecomdata.csv
/user/hive/warehouse/ecommerce/
```

Step 3: Start Hive

- 3.1. Open a terminal and run Hive:

```
hive
```

Step 4: Create a Hive Database

- 4.1. Inside the Hive shell, create a new database:
CREATE DATABASE ecommerce;

Step 5: Define External Tables:

```
USE ecommerce;
CREATE EXTERNAL TABLE ecommerce_data (
InvoiceNo STRING,
CompanyName STRING,
StockCode STRING,
Quantity INT,
InvoiceDate STRING,
UnitPrice DOUBLE,
CustomerID STRING,
Country STRING
)
ROW FORMAT DELIMITED FIELDS TERMINATED BY ','
LOCATION '/user/hive/warehouse/ecommerce/';
```

Step 6: Data Exploration and Understanding

6.1. Start exploring your data using Hive commands.

For example, to see the first five rows:

```
SELECT * FROM ecommerce_data LIMIT 5;
```

Step 7: Write Basic Queries

7.1. Begin by writing basic Hive queries.

For instance, to find the total number of records in your dataset:

```
SELECT COUNT(*) FROM ecommerce_data;
```

Step 8: Data Cleaning and Transformation

8.1. Data cleaning and transformation steps can be dataset-specific

-- To remove duplicate entries

```
INSERT OVERWRITE TABLE ecommerce_data  
SELECT DISTINCT *  
FROM ecommerce_data;
```

Step 9: Data Loading

9.1. Load data from the external table into an internal table.

```
CREATE TABLE customer AS  
SELECT  
customerid AS customer_id,  
companyname AS company_name,  
stockcode AS stock_code,  
quantity AS purchase_quantity,  
invoicedate AS invoice_date,  
unitprice AS unit_price,  
country AS customer_country  
FROM ecommerce_data;
```

Step 10: More Advanced Analysis and Partitioning

10.1. Calculate the total revenue or find the top-selling products.

```
USE ecommerce;  
-- Calculate total revenue  
SELECT SUM(price * quantity) AS total_revenue  
FROM ecommerce_data;
```

10.2. An example of partitioning by the `order_date` column:

– Create Table partitioned

```
CREATE TABLE ecommerce_partitioned (  
    InvoiceNo STRING,  
    CompanyName STRING,  
    StockCode STRING,  
    Quantity INT,  
    UnitPrice INT,  
    CustomerID STRING,  
    Country STRING  
)  
PARTITIONED BY (order_date DATE)  
STORED AS ORC  
LOCATION '/user/hive/warehouse/ecommerce_partitioned/';
```

– Load Data into the partitioned table

```
INSERT OVERWRITE TABLE ecommerce_partitioned PARTITION (order_date)  
SELECT  
    InvoiceNo,  
    CompanyName,  
    StockCode,  
    Quantity,  
    UnitPrice,  
    CustomerID,  
    Country,  
    CAST(InvoiceDate AS DATE) AS order_date  
FROM ecommerce_data;
```

10.3. Select All Data:

-- To retrieve all rows from the table, you can use the following query:

```
SELECT * FROM ecommerce_data;
```

10.4. Total Number of Records:

-- To find out how many records are in the table, you can use the COUNT function:

```
SELECT COUNT(*) AS record_count FROM ecommerce_data;
```

10.5. Filter Data by Date Range:

-- If you have a date column (e.g., InvoiceDate) and you want a specific date range:

```
SELECT *  
FROM ecommerce_data
```

```
WHERE InvoiceDate BETWEEN '2023-09-01' AND '2023-10-01';
```

10.6. Aggregations:

-- You can perform various aggregations on numeric columns.

For example, to find the total revenue:

```
SELECT SUM(UnitPrice * Quantity) AS total_revenue  
FROM ecommerce_data;
```

10.7. Average Unit Price:

-- To calculate the average unit price of products in the dataset:

```
SELECT AVG(UnitPrice) AS avg_unit_price  
FROM ecommerce_data;
```

10.8. Filter by Customer ID:

-- To filter data for a specific customer based on their CustomerID:

```
SELECT *  
FROM ecommerce_data  
WHERE CustomerID = '12345';
```